


Machine Learning

Probability




Reading: Bishop: Chap 1,2

Probability in Machine Learning

- Machine Learning tasks involve reasoning under uncertainty

Sources of uncertainty/randomness:

- Noise – variability in sensor measurements, partial observability, incorrect labels
- Finite sample size - Training and test data are randomly drawn instances



Hand-written digit recognition


Probability quantifies uncertainty!

Basic Probability Concepts

- Conceptual or physical, **repeatable** experiment with **random** outcome at any trial



Roll of dice



Nucleotide present at a DNA site



Time-space position of an aircraft on a radar screen

- Sample space S* - set of all possible outcomes. (can be finite or infinite.)

$$S \equiv \{1, 2, 3, 4, 5, 6\}$$

$$S \equiv \{A, T, C, G\}$$

$$S \equiv \{0, R_{\max}\} \times \{0, 360^\circ\} \times \{0, +\infty\}$$

- Event A* - any subset of S :

See "2", "4" or "6" in a roll

observe a "G" at a site

UA007 in angular location $\{45^\circ - 60^\circ\}$

Definition

- **Classical:** Probability of an event A is the relative frequency (limiting ratio of number of occurrences of event A to the total number of trials)


$$P(A) = \lim_{N \rightarrow \infty} \frac{N_A}{N}$$

E.g. $P(\{1\}) = 1/6$ $P(\{2,4,6\}) = 1/2$



Sample space S


Its area is 1, $P(S) = 1$




$P(A)$ - area of the oval

Definition

- ***Axiomatic (Kolmogorov):*** Probability of an event A is a number assigned to this event such that
- **$0 \leq P(A) \leq 1$** all probabilities are between 0 and 1




Area of A can't be smaller than 0




Area of A can't be larger than 1

Definition

- ***Axiomatic (Kolmogorov):*** Probability of an event A is a number assigned to this event such that
- $0 \leq P(A) \leq 1$ all probabilities are between 0 and 1
- $P(\phi) = 0$ probability of no outcome is 0




Area of A can't be smaller than 0




Area of A can't be larger than 1

Definition

- ***Axiomatic (Kolmogorov):*** Probability of an event A is a number assigned to this event such that
- $0 \leq P(A) \leq 1$ all probabilities are between 0 and 1
- $P(\phi) = 0$ probability of no outcome is 0
- $P(S) = 1$ probability of some outcome is 1




Area of A can't be smaller than 0



Area of A can't be larger than 1

Definition

- ***Axiomatic (Kolmogorov):*** Probability of an event A is a number assigned to this event such that
- $0 \leq P(A) \leq 1$ all probabilities are between 0 and 1
- $P(\phi) = 0$ probability of no outcome is 0
- $P(S) = 1$ probability of some outcome is 1
- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
probability of union of two events



Area of $A \cup B$ = Area of A + Area of B – Area of $A \cap B$

Definition

- ***Axiomatic (Kolmogorov):*** Probability of an event A is a number assigned to this event such that
- $0 \leq P(A) \leq 1$ all probabilities are between 0 and 1
- $P(\phi) = 0$ probability of no outcome is 0
- $P(S) = 1$ probability of some outcome is 1
- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
probability of union of two events
- ***Probability space*** is a sample space equipped with an assignment $P(A)$ to every event $A \subset S$ such that P satisfies the Kolmogorov axioms.

Theorems from the Axioms


- $0 \leq P(A) \leq 1$
- $P(\phi) = 0$
- $P(S) = 1$
- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

$$P(\neg A) = 1 - P(A)$$

$$\text{Proof: } P(A \cup \neg A) = P(S) = 1$$

$$P(A \cap \neg A) = P(\phi) = 0$$

$$1 = P(A) + P(\neg A) - 0 \Rightarrow P(\neg A) = 1 - P(A)$$




Theorems from the Axioms

- $0 \leq P(A) \leq 1$
- $P(\phi) = 0$
- $P(S) = 1$
- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

$$P(A) = P(A \cap B) + P(A \cap \neg B)$$

Proof: $P(A) = P(A \cap S) = P(A \cap (B \cup \neg B)) = P((A \cap B) \cup (A \cap \neg B))$
 $= P(A \cap B) + P(A \cap \neg B) - P((A \cap B) \cap (A \cap \neg B))$
 $= P(A \cap B) + P(A \cap \neg B) - P(\phi)$
 $= P(A \cap B) + P(A \cap \neg B)$



Why use probability?


- There have been many other approaches to handle uncertainty:
 - Fuzzy logic
 - Qualitative reasoning (Qualitative physics)
- “Probability theory is nothing but common sense reduced to calculation”
 - — Pierre Laplace, 1812.
- Any scheme for combining uncertain information really should obey these axioms
 - Di Finetti 1931 - If you gamble based on “uncertain beliefs” that satisfy these axioms, then you can’t be exploited by an opponent



Random Variable

- A **random variable** is a function that associates a unique numerical value $X(\omega)$ with every outcome $\omega \in S$ of an experiment.

(The value of the r.v. will vary from trial to trial as the experiment is repeated)



$X(\omega)$

$$P(X < 2) = P(\{\omega: X(\omega) < 2\})$$

- **Discrete r.v.:**
 - The outcome of a coin-toss $H = 1, T = 0$ (Binary)
 - The outcome of a dice-roll 1-6
- **Continuous r.v.:**
 - The location of an aircraft
- **Univariate r.v.:**
 - The outcome of a dice-roll 1-6
- **Multi-variate r.v.:**
 - The time-space position of an aircraft on radar screen
 - $$X = \begin{pmatrix} R \\ \Theta \\ t \end{pmatrix}$$

Discrete Probability Distribution

- **In the discrete case**, a probability distribution P on S (and hence on the domain of X) is an assignment of a non-negative real number $P(s)$ to each $s \in S$ (or each valid value of x) such that

$$0 \leq P(X=x) \leq 1 \quad X - \text{random variable}$$

$$\sum_x P(X = x) = 1 \quad x - \text{value it takes}$$

- **E.g. Bernoulli distribution with parameter θ**

$$P(x) = \begin{cases} 1 - \theta & \text{for } x = 0 \\ \theta & \text{for } x = 1 \end{cases} \Rightarrow P(x) = \theta^x (1 - \theta)^{1-x}$$



Discrete Probability Distribution

- In the discrete case**, a probability distribution P on S (and hence on the domain of X) is an assignment of a non-negative real number $P(s)$ to each $s \in S$ (or each valid value of x) such that

$0 \leq P(X=x) \leq 1$ X – random variable

$$\sum_{\mathbf{x}} \mathbf{P}(\mathbf{X} = \mathbf{x}) = 1 \quad \mathbf{x} - \text{value it takes}$$

E.g. Multinomial distribution with parameters $\theta_1, \dots, \theta_k$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_K \end{bmatrix}, \quad \text{where } \sum_j x_j = n$$

$$P(x) = \frac{n!}{x_1! x_2! \dots x_K!} \theta_1^{x_1} \theta_2^{x_2} \dots \theta_K^{x_K}$$

"Children"	"Aristocratic Budgets"	"Children"	"Aristocratic Budgets"	"C"
CHILDREN	NESTOROL MILLION	CHILDREN	NESTOROL MILLION	C1
ONEA	PSALMIST'S AX	WOMEN	PSALMIST'S AX	W
ONEA	SHEWOLDS PROGRAM	PEOPLE	SHEWOLDS PROGRAM	PI
ILD	MISSEATION GET	CHILD	MISSEATION GET	CI
SARS	MONARCH-BILLION	YCARIS	MONARCH-BILLION	YI
CHILD	FINCH FEDERAL	FAMILIES	FINCH FEDERAL	F1
ONE	MISSEAC YEAR	WORK	MISSEAC YEAR	W
LENTS	DETRACTOR SPENDING	TAKEN IS	DETRACTOR SPENDING	IS
YS	ARMONIT FIVE	SAYS	ARMONIT FIVE	S1
JULY	PERNIGAT STATE	FAMILY	PERNIGAT STATE	F1
ELIABLE	YACHTMAN PLAN	YACHTMAN PLAN	YACHTMAN PLAN	Y1
RY	OSWALD MONEY	WEN	OSWALD MONEY	W1
RYCHIT	THIRTY-SIX PROGRAMS	PERCENT	THIRTY-SIX PROGRAMS	PI
AZE	ACQUISITION GOVERNMENT	CARIS	ACQUISITION GOVERNMENT	CI
PH	LOWMOT CONGRESS	LAIT	LOWMOT CONGRESS	LI

Continuous Prob. Distribution

- A **continuous random variable** X can assume any value in an interval on the real line or in a region in a high dimensional space
 - X usually corresponds to a real-valued measurements of some property, **e.g., length, position, ...**
 - It is not possible to talk about the probability of the random variable assuming a particular value --- **$P(X=x) = 0$**
 - Instead, we talk about the probability of the random variable assuming a value within a given interval, or half interval

$$P(X \in [x_1, x_2])$$

$$P(X < x) = P(X \in [-\infty, x])$$

Continuous Prob. Distribution

- The probability of the random variable assuming a value within some given interval from x_1 to x_2 is defined to be the area under the graph of the probability density function between x_1 and x_2 .

- Probability mass: $P(X \in [x_1, x_2]) = \int_{x_1}^{x_2} p(x) dx$,

note that $\int_{-\infty}^{+\infty} p(x) dx = 1$.


- Cumulative distribution function (CDF):

$$F(x) = P(X \leq x) = \int_{-\infty}^x p(x') dx'$$

- Probability density function (PDF):

$$p(x) = \frac{d}{dx} F(x)$$

$$\int_{-\infty}^{+\infty} p(x) dx = 1; \quad p(x) \geq 0, \forall x$$



Car flow on Liberty Bridge (cooked up!)

What is the intuitive meaning of $p(x)$

- If

$$p(\mathbf{x}_1) = a \text{ and } p(\mathbf{x}_2) = b,$$

then when a value X is sampled from the distribution with density $p(x)$, you are a/b times as likely to find that X is “very close to” x than that \mathbf{x}_1 is “very close to” \mathbf{x}_2 .


- That is:

$$\lim_{h \rightarrow 0} \frac{P(x_1 - h < X < x_1 + h)}{P(x_2 - h < X < x_2 + h)} = \lim_{h \rightarrow 0} \frac{\int_{x_1-h}^{x_1+h} p(x) dx}{\int_{x_2-h}^{x_2+h} p(x) dx} \approx \frac{p(x_1) \times 2h}{p(x_2) \times 2h} = a/b$$

Continuous Distributions


- Uniform Probability Density Function

$$p(x) = 1/(b-a) \quad \text{for } a \leq x \leq b$$
$$= 0 \quad \text{elsewhere}$$



- Normal (Gaussian) Probability Density Function


$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/2\sigma^2}$$



- The distribution is **symmetric**, and is often illustrated as a **bell-shaped curve**.
- Two parameters, **μ (mean)** and **σ (standard deviation)**, determine the location and shape of the distribution.

- Exponential Probability Distribution

density: $p(x) = \frac{1}{\mu} e^{-x/\mu}$, CDF: $P(x \leq x_0) = 1 - e^{-x_0/\mu}$



Statistical Characterizations

- **Expectation:** the centre of mass, mean value, first moment

$$E(X) = \begin{cases} \sum_x xp(x) & \text{discrete} \\ \int_{-\infty}^{\infty} xp(x)dx & \text{continuous} \end{cases}$$

- **Variance:** the spread

$$\text{Var}(X) = \begin{cases} \sum_x (x - E(X))^2 p(x) & \text{discrete} \\ \int_{-\infty}^{\infty} (x - E(X))^2 p(x) dx & \text{continuous} \end{cases}$$

$$\text{Var}(X) = E(\underbrace{x - EX}_{\text{deviation}})^2 = E(x^2) - (EX)^2$$

Gaussian (Normal) density in 1D

- If $X \sim N(\mu, \sigma^2)$, the probability density function (pdf) of X is defined as

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/2\sigma^2}$$

$$E(X) = \mu$$


$$\text{var}(X) = \sigma^2$$

- Here is how we plot the pdf in matlab


```
xs=-3:0.01:3;
```

```
plot(xs,normpdf(xs,mu,sigma))
```

**Zero mean
Large variance**



**Zero mean
Small variance**



Note that a density evaluated at a point can be bigger than 1!



$$\int_{-\infty}^{\infty} p(x) dx = 1$$

Gaussian CDF 累积分布函数

- If $Z \sim N(0, 1)$, the cumulative density function is defined as

$$\begin{aligned}\Phi(x) &= \int_{-\infty}^x p(z) dz \\ &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x e^{-z^2/2} dz\end{aligned}$$

- This has no closed form expression, but is built in to most software packages (eg. `normcdf` in matlab stats toolbox).






Central limit theorem

- If (X_1, X_2, \dots, X_n) are i.i.d. (independent and identically distributed – to be covered next) random variables

- Then define

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$


- As $n \rightarrow \text{infinity}$,
- $p(\bar{X}) \rightarrow$ Gaussian with mean $E[X_i]$ and variance $\text{Var}[X_i]/n$



- Somewhat of a justification for assuming Gaussian distribution

Independence

- Training and test samples typically assumed to be i.i.d. (independent and identically distributed)



- A and B are **independent events** if

$$P(A \cap B) = P(A) * P(B)$$

- Outcome of A has no effect on the outcome of B (and vice versa).

E.g. Roll of two die
 $P(\{1\}, \{3\}) = 1/6 * 1/6 = 1/36$



Independence



- A, B and C are **pairwise independent** events if

$$P(A \cap B) = P(A) * P(B)$$

$$P(A \cap C) = P(A) * P(C)$$

$$P(B \cap C) = P(B) * P(C)$$

- A, B and C are **mutually independent** events if, in addition to pairwise independence,

$$P(A \cap B \cap C) = P(A) * P(B) * P(C)$$

Conditional Probability

- $P(A|B)$ = Probability of event A conditioned on event B having occurred


If $P(B) > 0$, then $P(A|B) = \frac{P(A \cap B)}{P(B)}$

E.g. H = "having a headache"

F = "coming down with Flu"

- $P(H) = 1/10$
- $P(F) = 1/40$
- $P(H|F) = 1/2$

Fraction of people with flu
that have a headache



- Corollary: **The Chain Rule**
- $P(A \cap B) = P(A|B) P(B)$

If A and B are independent, $P(A|B) = P(A)$

Conditional Independence

- A and B are **independent** if

$$P(A \cap B) = P(A) * P(B) \quad \equiv \quad P(A|B) = P(A)$$

- Outcome of B has no effect on the outcome of A (and vice versa).

- A and B are **conditionally independent given C** if


$$\underline{P(A \cap B|C) = P(A|C) * P(B|C)} \quad \equiv \quad \underline{P(A|B,C) = P(A|C)}$$

- Outcome of B has no effect on the outcome of A (and vice versa) if C is true.

Prior and Posterior Distribution

- Suppose that our propositions have a "causal flow"

e.g.,



- Prior or unconditional probabilities of propositions

e.g., $P(\text{Flu}) = 0.025$ and $P(\text{DrinkBeer}) = 0.2$

correspond to belief prior to arrival of any (new) evidence

- Posterior or conditional probabilities of propositions

e.g., $P(\text{Headache}|\text{Flu}) = 0.5$ and $P(\text{Headache}|\text{Flu}, \text{DrinkBeer}) = 0.7$

correspond to updated belief after arrival of new evidence

- Not always useful: $P(\text{Headache}|\text{Flu}, \text{Steelers win}) = 0.5$

Probabilistic Inference

H = "having a headache"

F = "coming down with Flu"

○ $P(H)=1/10$

○ $P(F)=1/40$

○ $P(H|F)=1/2$

- One day you wake up with a headache. You come with the following reasoning: "since 50% of flues are associated with headaches, so I must have a 50-50 chance of coming down with flu"

Is this reasoning correct?

Probabilistic Inference

H = "having a headache"

F = "coming down with Flu"


○ $P(H)=1/10$

○ $P(F)=1/40$

○ $P(H|F)=1/2$

● The Problem:

$$P(F|H) = ?$$



Probabilistic Inference


H = "having a headache"

F = "coming down with Flu"

- $P(H)=1/10$
- $P(F)=1/40$
- $P(H|F)=1/2$

● The Problem:

$$\begin{aligned} P(F|H) &= \frac{P(F \cap H)}{P(H)} \\ &= \frac{P(H|F)P(F)}{P(H)} \\ &= 1/8 \neq P(H|F) \end{aligned}$$




The Bayes Rule

- What we have just did leads to the following general expression:

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

This is Bayes Rule

Bayes, Thomas (1763) An essay towards solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society of London*, 53:370-418



Quiz

$$P(H)=1/10$$

$$P(F)=1/40$$

$$P(H|F)=1/2$$

$$P(F|H) = 1/8$$

● Which of the following statement is true?

$$P(F| \neg H) = 1 - P(F|H) \quad \times$$

$$P(\neg F|H) = 1 - P(F|H) \quad \checkmark$$

$$P(F| \neg H) = \frac{P(\neg H|F) P(F)}{P(\neg H)} = \frac{(1 - P(H|F)) P(F)}{1 - P(H)}$$

More General Forms of Bayes Rule

- $P(A|B) = \frac{P(B|A)P(A)}{P(B)}$

- Law of total probability 全概率

$$\begin{aligned} P(B) &= P(B \cap A) + P(B \cap \neg A) \\ &= P(B|A) P(A) + P(B|\neg A) P(\neg A) \end{aligned}$$


- $P(A|B) = \frac{P(B|A)P(A)}{P(B|A)P(A) + P(B|\neg A)P(\neg A)}$

More General Forms of Bayes Rule


- $$P(Y = y|X) = \frac{P(X|Y=y)P(Y=y)}{\sum_y P(X|Y = y)P(Y=y)}$$

$$P(Y|X \wedge Z) = \frac{P(X|Y \wedge Z)p(Y \wedge Z)}{P(X \wedge Z)} = \frac{P(X|Y \wedge Z)p(Y \wedge Z)}{P(X|\neg Y \wedge Z)p(\neg Y \wedge Z) + P(X|Y \wedge Z)p(\neg Y \wedge Z)}$$

E.g. $P(\text{Flu} | \text{Headhead} \wedge \text{DrankBeer})$



$$P(X \wedge Y \wedge Z) = P(Y \wedge Z)P(X|Y \wedge Z)$$



联合分布 Joint and Marginal Probabilities

- A **joint probability distribution** for a set of RVs (say X_1, X_2, X_3) gives the probability of every atomic event $P(X_1, X_2, X_3)$

○ $P(\text{Flu}, \text{DrinkBeer})$ = a 2×2 matrix of values:


	B	$\neg B$
F	0.005	0.02
$\neg F$	0.195	0.78

- $P(\text{Flu}, \text{DrinkBeer}, \text{Headache}) = ?$
- Every question about a domain can be answered by the joint distribution, as we will see later.
- A **marginal probability distribution** is the probability of every value that a single RV can take $P(X_1)$ $P(\text{Flu}) = ?$

Inference by enumeration

- **Start with a Joint Distribution**
- **Building a Joint Distribution of $M=3$ variables**
 - Make a truth table listing all combinations of values of your variables (if there are M Boolean variables then the table will have 2^M rows)
 - For each combination of values, say how probable it is.
 - Normalized, i.e., sums to 1

F	B	H	Prob
0	0	0	0.4
0	0	1	0.1
0	1	0	0.17
0	1	1	0.2
1	0	0	0.05
1	0	1	0.05
1	1	0	0.015
1	1	1	0.015




Inference with the Joint

- Once you have the JD you can ask for the probability of any atomic event consistent with your query

$$P(E) = \sum_{i \in E} P(row_i)$$

E.g. $E = \{(\neg F, \neg B, H), (\neg F, B, H)\}$

$\neg F$	$\neg B$	$\neg H$	0.4	
$\neg F$	$\neg B$	H	0.1	
$\neg F$	B	$\neg H$	0.17	
$\neg F$	B	H	0.2	
F	$\neg B$	$\neg H$	0.05	
F	$\neg B$	H	0.05	
F	B	$\neg H$	0.015	
F	B	H	0.015	



Inference with the Joint


- Compute Marginals

$P(\text{Flu} \wedge \text{Headache})$

$$= P(F \wedge H \wedge B) + P(F \wedge H \wedge \neg B)$$

Recall: Law of Total Probability


$\neg F$	$\neg B$	$\neg H$	0.4	
$\neg F$	$\neg B$	H	0.1	
$\neg F$	B	$\neg H$	0.17	
$\neg F$	B	H	0.2	
F	$\neg B$	$\neg H$	0.05	
F	$\neg B$	H	0.05	
F	B	$\neg H$	0.015	
F	B	H	0.015	



Inference with the Joint

- Compute Marginals
 $P(\text{Headache})$
 $= P(H \wedge F) + P(H \wedge \neg F)$
 $= P(H \wedge F \wedge B) + P(H \wedge F \wedge \neg B)$
 $+ P(H \wedge \neg F \wedge B) + P(H \wedge \neg F \wedge \neg B)$

$\neg F$	$\neg B$	$\neg H$	0.4	
$\neg F$	$\neg B$	H	0.1	
$\neg F$	B	$\neg H$	0.17	
$\neg F$	B	H	0.2	
F	$\neg B$	$\neg H$	0.05	
F	$\neg B$	H	0.05	
F	B	$\neg H$	0.015	
F	B	H	0.015	




Inference with the Joint

- Compute Conditionals

$$P(E_1|E_2) = \frac{P(E_1 \wedge E_2)}{P(E_2)}$$

$$= \frac{\sum_{i \in E_1 \cap E_2} P(\text{row}_i)}{\sum_{i \in E_2} P(\text{row}_i)}$$


$\neg F$	$\neg B$	$\neg H$	0.4	
$\neg F$	$\neg B$	H	0.1	
$\neg F$	B	$\neg H$	0.17	
$\neg F$	B	H	0.2	
F	$\neg B$	$\neg H$	0.05	
F	$\neg B$	H	0.05	
F	B	$\neg H$	0.015	
F	B	H	0.015	



Inference with the Joint

- Compute Conditionals
- $P(Flu|Headache) = \frac{P(Flu \wedge Headache)}{P(Headache)} = ?$
- General idea: Compute distribution on query variable by **fixing** evidence variables and **summing over** hidden variables

$\neg F$	$\neg B$	$\neg H$	0.4	
$\neg F$	$\neg B$	H	0.1	
$\neg F$	B	$\neg H$	0.17	
$\neg F$	B	H	0.2	
F	$\neg B$	$\neg H$	0.05	
F	$\neg B$	H	0.05	
F	B	$\neg H$	0.015	
F	B	H	0.015	



Where do probability distributions come from? 联合概率率

- Idea One: Human, Domain Experts
- Idea Two: Simpler probability facts and some algebra

e.g., $P(F)$

$P(B)$

$P(H|\neg F, B)$

$P(H|F, \neg B)$



$\neg F$	$\neg B$	$\neg H$	0.4	
$\neg F$	$\neg B$	H	0.1	
$\neg F$	B	$\neg H$	0.17	
$\neg F$	B	H	0.2	
F	$\neg B$	$\neg H$	0.05	
F	$\neg B$	H	0.05	
F	B	$\neg H$	0.015	
F	B	H	0.015	

- Use chain rule and independence assumptions to compute joint distribution


Where do probability distributions come from?

Four circles are positioned horizontally behind the title. From left to right: a solid light purple circle, a white circle with a light purple outline, a solid light purple circle, and a white circle with a light purple outline.

- Idea Three: Learn them from data!
 - A good chunk of this course is essentially about various ways of learning various forms of them!

Density Estimation 密度估计

- A Density Estimator learns a mapping from a set of attributes to a Probability



- Often know as parameter estimation if the distribution form is specified
 - Binomial, Gaussian...
- Some important issues:
 - Nature of the data (**iid, correlated, ...**)
 - Objective function (**MLE, MAP, ...**)
 - Algorithm (**simple algebra, gradient methods, EM, ...**)
 - Evaluation scheme (**likelihood on test data, predictability, consistency,**)

Parameter Learning from iid data

参数估计

- Goal: estimate distribution parameters θ from a dataset of independent, identically distributed (iid), fully observed, training cases

$$D = \{x_1, \dots, x_N\}$$

最大似然估计

- Maximum likelihood estimation (MLE)

1. One of the most common estimators
2. With iid and full-observability assumption, write $L(\theta)$ as the likelihood of the data:

$$\begin{aligned} L(\theta) &= P(D; \theta) = P(x_1, x_2, \dots, x_N; \theta) \\ &= P(X_1; \theta) P(X_2; \theta) \dots P(X_N; \theta) \\ &= \prod_i^N P(X_i; \theta) \end{aligned}$$

3. pick the setting of parameters most likely to have generated the data we saw:

$$\hat{\theta} = \underset{\theta}{\operatorname{argmax}} L(\theta) = \underset{\theta}{\operatorname{argmax}} \log(L(\theta))$$

MLE: 模型已定, 参数未知 (是个固定值)

Example 1: Bernoulli model

- **Data:** 目的为找出参数, 使观测数据出现的可能
○ We observed N iid coin tossing: $D = \{1, 0, 1, \dots, 0\}$ 性最大
- **Model:** 结果准确性依赖于潜在数据分布

$$P(x) = \begin{cases} 1-\theta & \text{for } x=0 \\ 0 & \text{for } x=1 \end{cases} \Rightarrow P(x) = \theta^x (1-\theta)^{1-x}$$

- **How to write the likelihood of a single observation x_i ?**

$$P(x_i) = \theta^{x_i} (1-\theta)^{1-x_i}$$

- **The likelihood of dataset $D = \{x_1, \dots, x_N\}$:**

$$\begin{aligned} L(\theta) &= P(x_1, x_2, \dots, x_N; \theta) = \prod_{i=1}^N P(x_i; \theta) = \prod_{i=1}^N (\theta^{x_i} (1-\theta)^{1-x_i}) \\ &= \theta^{\sum_{i=1}^N x_i} (1-\theta)^{\sum_{i=1}^N 1-x_i} = \theta^{\# \text{head}} (1-\theta)^{\# \text{tails}} \end{aligned}$$

$$L(\theta) = \theta^{\sum x_i} (1-\theta)^{n - \sum x_i}$$

MLE

$$\begin{aligned}\ln L(\theta) &= \ln \theta^{\sum x_i} + \ln (1-\theta)^{n-\sum x_i} \\ &= \sum x_i \ln \theta + (n - \sum x_i) \ln (1-\theta) \\ \frac{\partial \ln L(\theta)}{\partial \theta} &= \frac{\sum x_i}{\theta} - \frac{n - \sum x_i}{1-\theta} = 0 \\ \theta &= \frac{\sum x_i}{n}\end{aligned}$$

- Objective function:

$$\ell(\theta) = \log L(\theta) = \log \theta^{n_h} (1-\theta)^{n_t} = n_h \log \theta + (N - n_h) \log(1-\theta)$$

- We need to maximize this w.r.t. θ

- Take derivatives w.r.t θ

$$\frac{\partial \ell}{\partial \theta} = \frac{n_h}{\theta} - \frac{N - n_h}{1-\theta} = 0$$



$$\hat{\theta}_{MLE} = \frac{n_h}{N}$$

or $\hat{\theta}_{MLE} = \frac{1}{N} \sum_i x_i$

Frequency as sample mean

$$n_h - \theta n_h = \theta N - \theta n_h$$

$$\theta = \frac{n_h}{N}$$

- Sufficient statistics

The counts, n_h , where $n_h = \sum_i x_i$, are sufficient statistics of data D

$\sum_{i=1}^N x_i$
充分统计量

Example 2: univariate normal

单变量

- Data:

- We observed N iid real samples:
 $D = \{-0.1, 10, 1, -5.2, \dots, 3\}$

$\ell(\theta)$

- Model: $P(x) = (2\pi\sigma^2)^{-1/2} \exp\{-(x - \mu)^2 / 2\sigma^2\}$ $\theta = (\mu, \sigma^2)$

- Log likelihood:

$$\ell(\theta) = \log L(\theta) = \prod_{i=1}^N P(x_i) = -\frac{N}{2} \log(2\pi\sigma^2) - \frac{1}{2} \sum_{i=1}^N \frac{(x_i - \mu)^2}{\sigma^2}$$

- MLE: take derivative and set to zero:

$$\frac{\partial \ell}{\partial \mu} = (1/\sigma^2) \sum_n (x_n - \mu)$$

$$\frac{\partial \ell}{\partial \sigma^2} = -\frac{N}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_n (x_n - \mu)^2$$






$$\mu_{\text{MLE}} = \frac{1}{N} \sum_n x_n$$

$$\sigma_{\text{MLE}}^2 = \frac{1}{N} \sum_n (x_n - \mu_{\text{ML}})^2$$

Overfitting

过拟合



- Recall that for Bernoulli Distribution, we have

$$\hat{\theta}_{ML}^{head} = \frac{n^{head}}{n^{head} + n^{tail}}$$

- What if we tossed too few times so that we saw zero head? We have $\hat{\theta}_{ML}^{head} = 0$, and we will predict that the probability of seeing a head next is zero!!!

- The rescue “smoothing”:

- Where n' is known as the pseudo- (imaginary) count

$$\hat{\theta}_{ML}^{head} = \frac{n^{head} + n'}{n^{head} + n^{tail} + n'}$$

- But can we make this more formal? *MAP*

Bayesian Learning

- The Bayesian Rule:

$$P(\theta | \mathcal{D}) = \frac{P(\mathcal{D} | \theta)P(\theta)}{P(\mathcal{D})}$$

即 $P(\mathcal{D} | \theta)$

似然函数

观测数据前, 有些关于参数的假设, 以先验概率的形式

常数, 归一化

- Or equivalently,

$$\underbrace{P(\theta | \mathcal{D})}_{\text{posterior}} \propto \underbrace{P(\mathcal{D} | \theta)}_{\text{likelihood}} \underbrace{P(\theta)}_{\text{prior}}$$

正比于

(Belief about coin toss probability)

最大后验概率


- MAP estimate: $\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} P(\theta | \mathcal{D})$

- If prior is uniform, MLE = MAP

Bayesian estimation for Bernoulli

- **Beta(α, β) distribution:**

$$P(\theta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1} (1 - \theta)^{\beta-1} = B(\alpha, \beta) \theta^{\alpha-1} (1 - \theta)^{\beta-1}$$



- **Posterior distribution of θ :**

$$P(\theta | D) = \frac{p(x_1, \dots, x_N | \theta) p(\theta)}{p(x_1, \dots, x_N)} \propto \theta^{n_h} (1 - \theta)^{n_t} \times \theta^{\alpha-1} (1 - \theta)^{\beta-1} = \theta^{n_h + \alpha - 1} (1 - \theta)^{n_t + \beta - 1}$$

Beta($\alpha + n_h, \beta + n_t$)

- Notice the **isomorphism** of the **posterior to the prior**,
- such a prior is called a **conjugate prior**
- α and β are **hyperparameters** (parameters of the prior) and correspond to the number of “virtual” heads/tails (pseudo counts)

MAP

- Posterior distribution of θ :

$$P(\theta | x_1, \dots, x_N) = \frac{p(x_1, \dots, x_N | \theta) p(\theta)}{p(x_1, \dots, x_N)} \propto \theta^{n_h} (1 - \theta)^{n_t} \times \theta^{\alpha-1} (1 - \theta)^{\beta-1} = \theta^{n_h + \alpha - 1} (1 - \theta)^{n_t + \beta - 1}$$

- Maximum a posteriori (MAP) estimation:

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} \log P(\theta | x_1, \dots, x_N)$$

- Posterior mean estimation:

$$\hat{\theta}_{\text{MAP}} = \frac{n_h + \alpha}{N + \alpha + \beta}$$

Beta parameters
can be understood
as pseudo-counts

- With enough data, prior is forgotten

Dirichlet distribution

- number of heads in N flips of a two-sided coin
 - follows a binomial distribution
 - Beta is a good prior (conjugate prior for binomial)
- what it's not two-sided, but k-sided?
 - follows a multinomial distribution
 - Dirichlet distribution is the conjugate prior

$$P(\theta_1, \theta_2, \dots, \theta_K) = \frac{1}{B(\alpha)} \prod_i^K \theta_i^{(\alpha_i - 1)}$$

Lejeune Dirichlet



Johann Peter Gustav Lejeune Dirichlet

Born	13 February 1805 Düren, French Empire
Died	5 May 1859 (aged 54) Göttingen, Hanover
Residence	 Germany
Nationality	 German
Fields	Mathematician
Institutions	University of Berlin University of Breslau University of Göttingen
Alma mater	University of Bonn
Doctoral advisor	Simeon Poisson Joseph Fourier
Doctoral students	Ferdinand Eisenstein Leopold Kronecker Rudolf Lipschitz Carl Wilhelm Borchardt
Known for	Dirichlet function Dirichlet eta function

Estimating the parameters of a distribution

- Maximum Likelihood estimation (MLE) Choose value that maximizes the probability of observed data

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} P(D | \theta)$$

- Maximum a posteriori (MAP) estimation Choose value that is most probable given observed data and prior belief

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} P(\theta | D) = \arg \max_{\theta} P(D | \theta) P(\theta)$$

先验概率

MLE vs MAP (Frequentist vs Bayesian)

- Frequentist/MLE approach:
 - θ is unknown constant, estimate from data
- Bayesian/MAP approach: 常量
 - θ is a random variable, assume a probability distribution 变量
- Drawbacks
 - **MLE**: Overfits if dataset is too small
 - **MAP**: Two people with different priors will end up with different estimates

Bayesian estimation for normal distribution

- Normal Prior:

$$P(\mu) = (2\pi\tau^2)^{-1/2} \exp\left\{-\frac{(\mu - \mu_0)^2}{2\tau^2}\right\}$$

- Joint probability:

$$P(\mathbf{x}, \mu) = (2\pi\sigma^2)^{-N/2} \exp\left\{-\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2\right\} \\ \times (2\pi\tau^2)^{-1/2} \exp\left\{-\frac{(\mu - \mu_0)^2}{2\tau^2}\right\}$$

- Posterior:

$$P(\mu | \mathbf{x}) = (2\pi\tilde{\sigma}^2)^{-1/2} \exp\left\{-\frac{(\mu - \tilde{\mu})^2}{2\tilde{\sigma}^2}\right\}$$

where $\tilde{\mu} = \frac{N/\sigma^2}{N/\sigma^2 + 1/\tau^2} \bar{x} + \frac{1/\tau^2}{N/\sigma^2 + 1/\tau^2} \mu_0$, and $\tilde{\sigma}^2 = \left(\frac{N}{\sigma^2} + \frac{1}{\tau^2}\right)^{-1}$

 Sample mean

Probability Review

What you should know:

- Probability basics
 - random variables, events, sample space, conditional probs, ...
 - independence of random variables
 - Bayes rule
 - Joint probability distributions
 - calculating probabilities from the joint distribution
- Point estimation
 - maximum likelihood estimates
 - maximum a posteriori estimates
 - distributions – binomial, Beta, Dirichlet, ...